

Case Study 2 AKSTA Statistical Computing

Hanna Kienast

Iulia Mihaela Enache

Kateryna Ponomarenko

2024-04-30

Task

a

Load in R the following data sets which you can find in TUWEL. For each data set, ensure that missing values are read in properly, that column names are unambiguous. Each data set should contain at the end only two columns: country and the variable.

```
library(readr)
```

```
## Warning: Paket 'readr' wurde unter R Version 4.3.3 erstellt
```

```
library(tidyr)
```

```
## Warning: Paket 'tidyr' wurde unter R Version 4.3.2 erstellt
```

```
library(dplyr)
```

```
## Warning: Paket 'dplyr' wurde unter R Version 4.3.2 erstellt
```

```
##
```

```
## Attache Paket: 'dplyr'
```

```
## Die folgenden Objekte sind maskiert von 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## Die folgenden Objekte sind maskiert von 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(readxl)
```

```
## Warning: Paket 'readxl' wurde unter R Version 4.3.3 erstellt
```

```
library(data.table)
```

```
## Warning: Paket 'data.table' wurde unter R Version 4.3.3 erstellt
```

```
##
```

```
## Attache Paket: 'data.table'
```

```
## Die folgenden Objekte sind maskiert von 'package:dplyr':
```

```
##
```

```
##      between, first, last
```

```
require(stringr)
```

```
## Lade nötiges Paket: stringr
```

```
median_age_data <- read.fwf("rawdata_343.txt", widths = c(8,66,4,18), skip=2)
names(median_age_data) <- c("Rank", "Country", "Median_Age")
median_age_data <- median_age_data[, c("Country", "Median_Age")]
head(median_age_data)
```

```
##                                     Country Median_Age
## 1 Monaco                                     55.4
## 2 Japan                                     48.6
## 3 Saint Pierre and Miquelon                48.5
## 4 Germany                                    47.8
## 5 Italy                                      46.5
## 6 Andorra                                    46.2
```

```
migration_data <- read.fwf("rawdata_347.txt", widths = c(8,64,6,18), skip=2)
names(migration_data) <- c("Rank", "Country", "Migration_rate", "Date")
migration_data <- migration_data[, c("Country", "Migration_rate")]
head(migration_data)
```

```
##                                     Country
## 1 Syria
## 2 British Virgin Islands
## 3 Luxembourg
## 4 Cayman Islands
## 5 Singapore
## 6 Anguilla
## Migration_rate
## 1          27.1
## 2          15.5
## 3          13.3
## 4          13.0
## 5          11.8
## 6          11.1
```

```
unemployment_data <- read.csv("rawdata_373.csv")
names(unemployment_data) <- c("Country", "Youth_unemployment_rate")
head(unemployment_data)
```

```
##                                     Country
## 1 French Polynesia
## 2 Kosovo
## 3 South Africa
## 4 Libya
## 5 Eswatini
## 6 Saint Lucia
## Youth_unemployment_rate
## 1          56.7
## 2          55.4
## 3          53.4
## 4          48.7
## 5          47.1
## 6          46.2
```

b

Merge the data sets containing raw data using dplyr function on the unique keys. Keep the union of all observations in the tables. What key are you using for merging? Return the dimension of the merged data set. We use the column “country” for merging. We encountered an issue where the column for migration rate was always NaN so we had to trim the whitespaces.

```
median_age_data$Country <- trimws(median_age_data$Country)
migration_data$Country <- trimws(migration_data$Country)
unemployment_data$Country <- trimws(unemployment_data$Country)

joined <- left_join(median_age_data, unemployment_data, by = join_by(Country))
joined_data <- left_join(joined, migration_data, by = join_by(Country == Country))
head(joined_data)
```

```
##          Country Median_Age Youth_unemployment_rate Migration_rate
## 1 Monaco      55.4      26.6              8.3
## 2 Japan       48.6       3.6              0.0
## 3 Saint Pierre and Miquelon 48.5      NA             -7.7
## 4 Germany     47.8       6.2              1.5
## 5 Italy       46.5      32.2              3.2
## 6 Andorra     46.2      NA              0.0
```

```
dim(joined_data)
```

```
## [1] 227  4
```

c

You will acquire more country level information such as the classification of the country based on income. Such an information can be found at <https://datahelpdesk.worldbank.org/knowledgebase/articles/906519>. From there extract the classification for 2020 into low/lower-middle/upper-middle/high income countries

```
income_data <- read_excel("OGHIST.xlsx", sheet="Country Analytical History", col_names = TRUE)
```

```
## New names:
```

```
## * '' -> '...1'  
## * '' -> '...3'  
## * '' -> '...4'  
## * '' -> '...5'  
## * '' -> '...6'  
## * '' -> '...7'  
## * '' -> '...8'  
## * '' -> '...9'  
## * '' -> '...10'  
## * '' -> '...11'  
## * '' -> '...12'  
## * '' -> '...13'  
## * '' -> '...14'  
## * '' -> '...15'  
## * '' -> '...16'  
## * '' -> '...17'  
## * '' -> '...18'  
## * '' -> '...19'  
## * '' -> '...20'  
## * '' -> '...21'  
## * '' -> '...22'  
## * '' -> '...23'  
## * '' -> '...24'  
## * '' -> '...25'  
## * '' -> '...26'  
## * '' -> '...27'  
## * '' -> '...28'  
## * '' -> '...29'  
## * '' -> '...30'  
## * '' -> '...31'  
## * '' -> '...32'  
## * '' -> '...33'  
## * '' -> '...34'  
## * '' -> '...35'  
## * '' -> '...36'  
## * '' -> '...37'  
## * '' -> '...38'
```

```
income_data <- income_data[, c("World Bank Analytical Classifications", "...36")]  
income_data <- income_data[11:nrow(income_data), ]  
names(income_data) <- c("Country", "Income_class")  
head(income_data)
```

```
## # A tibble: 6 x 2  
##   Country      Income_class  
##   <chr>        <chr>  
## 1 Afghanistan L  
## 2 Albania     UM  
## 3 Algeria     LM
```

```
## 4 American Samoa UM
## 5 Andorra          H
## 6 Angola           LM
```

d

Merge this information to the data set in b. 1. What are the common variables? Can you merge using them? Why or why not?

We can merge the dataframes on the column Country.

```
data <- left_join(joined_data, income_data, by = join_by(Country == Country))
head(data)
```

```
##           Country Median_Age Youth_unemployment_rate Migration_rate
## 1           Monaco      55.4              26.6              8.3
## 2             Japan      48.6              3.6              0.0
## 3 Saint Pierre and Miquelon 48.5              NA             -7.7
## 4             Germany      47.8              6.2              1.5
## 5              Italy      46.5             32.2              3.2
## 6             Andorra      46.2              NA              0.0
##   Income_class
## 1            H
## 2            H
## 3          <NA>
## 4            H
## 5            H
## 6            H
```

2. A reliable merging for countries are ISO codes as they are standardized across data sources. Download the mapping of ISO codes to countries from <https://www.cia.gov/the-world-factbook/references/country-data-codes/> and load it into R.

```
country_code <- fread("Country Data Codes.csv", sep = ",", header = TRUE, fill=TRUE, stringsAsFactors =
names(country_code)[names(country_code) == 'ISO 3166'] <- 'ISO'
names(country_code)[names(country_code) == 'Name'] <- 'Country'

country_code <- country_code %>%
  mutate(
    Country = ifelse(
      str_detect(GENC, "[a-z]"),
      paste(Country, GENC, sep = ", "),
      Country
    ),
    GENC = ifelse(
      str_detect(GENC, "[a-z]"),
      NA,
      GENC
    )
  )

country_code <- country_code[, c("Country", "ISO")]
```

```
country_code <- country_code %>%
  mutate(across(everything(), ~ str_replace_all(., ' ', '')))

head(country_code)
```

```
##           Country      ISO
##           <char>    <char>
## 1:  Afghanistan AF|AFG|004
## 2:    Akrotiri    -
## 3:    Albania AL|ALB|008
## 4:    Algeria DZ|DZA|012
## 5: American Samoa AS|ASM|016
## 6:    Andorra AD|AND|020
```

3. Merge the data sets using the ISO codes.

```
country_code$Country <- trimws(country_code$Country)
country_code$ISO <- trimws(country_code$ISO)

data <- left_join(data, country_code, by = join_by(Country == Country))
names(data)[names(data) == 'ISO.y'] <- 'ISO'
data$ISO.x <- NULL
head(data)
```

```
##           Country Median_Age Youth_unemployment_rate Migration_rate
## 1             Monaco      55.4              26.6           8.3
## 2              Japan      48.6              3.6           0.0
## 3 Saint Pierre and Miquelon      48.5              NA          -7.7
## 4              Germany      47.8              6.2           1.5
## 5              Italy      46.5             32.2           3.2
## 6              Andorra      46.2              NA           0.0
## Income_class      ISO
## 1             H MC|MCO|492
## 2             H JP|JPN|392
## 3      <NA> PM|SPM|666
## 4             H DE|DEU|276
## 5             H IT|ITA|380
## 6             H AD|AND|020
```

e

Introduce into the data set information on continent for each country and subcontinent (region). You should find a way to gather this data. You can find an appropriate online resource, download the data and merge the information with the existing data set. Name the merged data set `df_vars`.

We found this dataset on kaggle: <https://www.kaggle.com/datasets/andradaolteanu/country-mapping-iso-continent-region?resource=download> However, despite ISO being standardized, we could not find a dataset where the ISO layout was the same as the one provided from the previous link. For making the merging work, we had to merge 3 columns to create a similar column to ISO from the previous dataframe.

```

continents <- read.csv("continents2.csv")

continents <- continents %>%
  mutate(
    country.code = str_pad(country.code, width = 3, pad = "0"),

    ISO = paste(alpha.2, alpha.3, country.code, sep = "|")
  )

names(continents)[names(continents) == 'name'] <- 'Country'
continents <- continents[, c("Country", "region", "sub.region", "ISO")]
head(continents)

```

```

##           Country region      sub.region      ISO
## 1  Afghanistan  Asia  Southern Asia AF|AFG|004
## 2   Åland Islands Europe Northern Europe AX|ALA|248
## 3      Albania  Europe Southern Europe AL|ALB|008
## 4      Algeria  Africa Northern Africa DZ|DZA|012
## 5 American Samoa Oceania      Polynesia AS|ASM|016
## 6      Andorra  Europe Southern Europe AD|AND|020

```

```

df_vars <- left_join(data, continents, by = join_by(Country == Country))
names(df_vars)[names(df_vars) == 'ISO.x'] <- 'ISO'
df_vars$ISO.y <- NULL
dim(df_vars)

```

```
## [1] 227 8
```

```
head(df_vars)
```

```

##           Country Median_Age Youth_unemployment_rate Migration_rate
## 1           Monaco      55.4              26.6              8.3
## 2            Japan      48.6              3.6              0.0
## 3 Saint Pierre and Miquelon      48.5              NA             -7.7
## 4            Germany      47.8              6.2              1.5
## 5             Italy      46.5             32.2              3.2
## 6      Andorra      46.2              NA              0.0
##  Income_class      ISO region      sub.region
## 1      H MC|MCO|492  Europe  Western Europe
## 2      H JP|JPN|392   Asia   Eastern Asia
## 3    <NA> PM|SPM|666 Americas Northern America
## 4      H DE|DEU|276  Europe  Western Europe
## 5      H IT|ITA|380  Europe  Southern Europe
## 6      H AD|AND|020  Europe  Southern Europe

```

f

Discuss on the tidyness of the data set `df_vars`. What are the observational units, what are the variables? What can be considered fixed vs measured variables? Tidy the data if needed.

The obtained data set `df_vars` is almost tidy, while it almost fully satisfies the requirements to the data to be considered tidy: our data set is the collection of quantitative and qualitative values and they are organized in a way, that each value belongs to an observation and a variable. Our data is organized in a way that each row represents an observational unit (a country) and each column is a variable. All our variables have the same unit and measure the same attribute.

In our case the fixed variables are `Country` (char), `region` (char), `sub.region` (char), `ISO` (char), `Income_class` (char). The left variables `Median_Age` (double), `Migration_rate` (double) and `Youth_unemployment_rate` (double) are measured variables in our case. Typically, the fixed variables are put in the beginning of the data set, while measured are put after them. So we will perform this small change to have our data perfectly tidy.

Also, since in the following tasks the region will be pointed to as “continent”, we decided to rename the name of this variable.

```
df_vars <- df_vars %>%
  select(ISO, Country, region, sub.region, Income_class, Median_Age, Youth_unemployment_rate, Migration_rate)

names(df_vars)[names(df_vars) == "region"] <- "continent"

head(df_vars)
```

```
##      ISO      Country continent      sub.region Income_class
## 1 MC|MCO|492      Monaco   Europe  Western Europe          H
## 2 JP|JPN|392       Japan    Asia    Eastern Asia          H
## 3 PM|SPM|666 Saint Pierre and Miquelon Americas Northern America <NA>
## 4 DE|DEU|276      Germany   Europe  Western Europe          H
## 5 IT|ITA|380       Italy    Europe  Southern Europe          H
## 6 AD|AND|020      Andorra   Europe  Southern Europe          H
## Median_Age Youth_unemployment_rate Migration_rate
## 1      55.4              26.6           8.3
## 2      48.6              3.6           0.0
## 3      48.5              NA          -7.7
## 4      47.8              6.2           1.5
## 5      46.5             32.2           3.2
## 6      46.2              NA           0.0
```

g

Make a frequency table for the status variable in the merged data set. Briefly comment on the results.

```
income_status_frequency <- df_vars %>%
  count(Income_class, name = "frequency") %>%
  mutate(
    percentage = 100 * (frequency / sum(frequency))
  )

print(income_status_frequency)
```

```
##      Income_class frequency percentage
## 1              H         66   29.07489
## 2              L         23   10.13216
```



```
## 3          LM          45    19.82379
## 4          UM          50    22.02643
## 5        <NA>         43    18.94273
```

From the generated table we can conclude that the high income countries are represented at most, while low income countries at least. The representation of lower-middle and upper-middle income countries are on the same level. Also, a lot of observations (about 19%) are missing the value for the variable `Income_class` in our data set.

h

What is the distribution of income status in the different continents? Compute the absolute frequencies as well as the relative frequency of status within each continent. Briefly comment on the results.

```
income_distribution_within_continent <- df_vars %>%
  group_by(continent, Income_class) %>%
  summarise(frequency = n(), .groups = 'drop')

income_distribution_within_continent <- income_distribution_within_continent %>%
  group_by(continent) %>%
  mutate(relative_frequency = frequency / sum(frequency))

print(income_distribution_within_continent)
```

```
## # A tibble: 26 x 4
## # Groups:   continent [6]
##   continent Income_class frequency relative_frequency
##   <chr>      <chr>          <int>          <dbl>
## 1 Africa    H              1            0.0204
## 2 Africa    L             20            0.408
## 3 Africa    LM             19            0.388
## 4 Africa    UM              7            0.143
## 5 Africa    <NA>           2            0.0408
## 6 Americas  H             13            0.302
## 7 Americas  LM              6            0.140
## 8 Americas  UM             17            0.395
## 9 Americas  <NA>           7            0.163
## 10 Asia     H             10            0.222
## # i 16 more rows
```

We can conclude that the low and lower-income classes are prevailing in Africa, while the biggest group in Europe and Oceania is high income class. In Asia the lower-middle income class is the biggest one followed by the upper-middle income class. In Americas the prevailing income class is upper-middle and in our data set there is no low income class countries represented. Also, it should be noted that each continent has some countries with no income class specification available (Africa $\approx 4\%$, Americas $\approx 16\%$, Asia $\approx 18\%$, Europe $\approx 13\%$ and Oceania 10%). In addition, some countries have values for variable `Income_class`, while no for continent.

i

From h. identify the countries which are the only ones in their respective group. Explain in few words the output.

```
unique_countries_in_their_group <- df_vars %>%
  group_by(continent, Income_class) %>%
  filter(n() == 1) %>%
  ungroup()

print(unique_countries_in_their_group)
```

```
## # A tibble: 3 x 8
##   ISO      Country    continent sub.region      Income_class Median_Age
##   <chr>    <chr>      <chr>    <chr>      <chr>         <dbl>
## 1 UA|UKR|804 Ukraine    Europe    Eastern Europe    LM          41.2
## 2 SC|SYC|690 Seychelles Africa     Sub-Saharan Africa H          36.8
## 3 AF|AFG|004 Afghanistan Asia       Southern Asia     L          19.5
## # i 2 more variables: Youth_unemployment_rate <dbl>, Migration_rate <dbl>
```

We can see that Ukraine is the only lower-middle income country in Europe, while the only high income country in Africa is Seychelles and the only low income country in Asia is Afghanistan.

j

For each continent count the number of sub-regions in the data set. How granular are the subcontinents that you employ in the analysis?

```
sub_region_counts_for_continents <- df_vars %>%
  group_by(continent) %>%
  summarise(number_of_sub_regions = n_distinct(sub.region), .groups = 'drop')

print(sub_region_counts_for_continents)
```

```
## # A tibble: 6 x 2
##   continent number_of_sub_regions
##   <chr>          <int>
## 1 Africa                2
## 2 Americas              2
## 3 Asia                  5
## 4 Europe                4
## 5 Oceania               4
## 6 <NA>                  1
```

It is seen that the Asia has the highest granularity, followed by Europe and Oceania, while Africa and Americas are represented only by two sub-regions in our data set.

k

Look at the frequency distribution of income status in the subregions of North- and South-Americas. Comment on the results.

```
income_distribution_americas <- df_vars %>%
  filter(continent %in% "Americas") %>%
  group_by(sub.region, Income_class) %>%
  summarise(frequency = n(), .groups = 'drop') %>%
  group_by(sub.region) %>%
  mutate(relative_frequency = frequency / sum(frequency))

print(income_distribution_americas)
```

```
## # A tibble: 6 x 4
## # Groups:   sub.region [2]
##   sub.region      Income_class frequency relative_frequency
##   <chr>          <chr>          <int>          <dbl>
## 1 Latin America and the Caribbean H             9           0.237
## 2 Latin America and the Caribbean LM             6           0.158
## 3 Latin America and the Caribbean UM            17           0.447
## 4 Latin America and the Caribbean <NA>           6           0.158
## 5 Northern America      H             4           0.8
## 6 Northern America      <NA>           1           0.2
```

We obtained that in the Northern America all countries are high income and also we have one country with no specified income status. In the Southern America prevailing number of countries are upper-middle class income countries, around 24% of countries are high income.

1.

Dig deeper into the low-middle income countries of the Americas. Which ones are they? Are they primarily small island states in the Caribbean? Comment.

```
low_middle_income_americas <- df_vars %>%
  filter(continent %in% "Americas" &
         Income_class == "LM")

print(low_middle_income_americas)
```

```
##      ISO      Country continent      sub.region Income_class
## 1 SV|SLV|222 El Salvador Americas Latin America and the Caribbean LM
## 2 NI|NIC|558 Nicaragua Americas Latin America and the Caribbean LM
## 3 BO|BOL|068 Bolivia Americas Latin America and the Caribbean LM
## 4 HN|HND|340 Honduras Americas Latin America and the Caribbean LM
## 5 HT|HTI|332 Haiti Americas Latin America and the Caribbean LM
## 6 BZ|BLZ|084 Belize Americas Latin America and the Caribbean LM
##   Median_Age Youth_unemployment_rate Migration_rate
## 1      27.7      9.6      -4.8
## 2      27.3      8.5      -2.4
## 3      25.3      6.9      -0.3
## 4      24.4     10.7      -1.4
## 5      24.1      NA      -1.9
## 6      23.9     15.3      -1.0
```

No, these countries are not island states (except Haiti), however they all except Bolivia have the coastal line. Also, they all are quite differ in terms of the size.

m

Create a table of average values for median age, youth unemployment rate and net migration rate separated into income status. Make sure that in the output, the ordering of the income classes is proper (i.e., L, LM, UM, H or the other way around). Briefly comment the results.

We will calculate it by ignoring the NA values. Also, we will exclude from the final data frame the values for the observations with no income class specified

```
table_of_average_values <- df_vars %>%
  filter(!is.na(Income_class)) %>%
  mutate(Income_class = factor(Income_class, levels = c("L", "LM", "UM", "H"))) %>%
  group_by(Income_class) %>%
  summarise(
    Avg_Median_Age = mean(Median_Age, na.rm = TRUE),
    Avg_Youth_Unemployment_Rate = mean(Youth_unemployment_rate, na.rm = TRUE),
    Avg_Migration_Rate = mean(Migration_rate, na.rm = TRUE),
    .groups = 'drop'
  )

print(table_of_average_values)
```

```
## # A tibble: 4 x 4
##   Income_class Avg_Median_Age Avg_Youth_Unemployment_Rate Avg_Migration_Rate
##   <fct>          <dbl>                <dbl>                <dbl>
## 1 L             18.4                  13.2                 -1.74
## 2 LM            24.9                  16.4                 -1.50
## 3 UM            31.5                  21.8                 -4.48
## 4 H             39.3                  16.7                  2.18
```

We can trace the growth of the average median age of the population with the drop of the status of country's income. In addition, we observe that the highest average unemployment youth rate is experienced by the upper-middle status income countries. The high and lower-middle income countries have the same almost the same average unemployment youth and the indicator for low income status countries is the smallest. Average migration rate in countries of all statuses except high income status is negative, being the most negative in countries with upper-middle income status.

n

Look also at the standard deviation instead of the mean in m. Do you gain additional insights? Briefly comment the results.

```
table_of_sd_values <- df_vars %>%
  filter(!is.na(Income_class)) %>%
  mutate(Income_class = factor(Income_class, levels = c("L", "LM", "UM", "H"))) %>%
  group_by(Income_class) %>%
  summarise(
    Sd_Median_Age = sd(Median_Age, na.rm = TRUE),
    Sd_Youth_Unemployment_Rate = sd(Youth_unemployment_rate, na.rm = TRUE),
    Sd_Migration_Rate = sd(Migration_rate, na.rm = TRUE),
    .groups = 'drop'
  )
```

```
print(table_of_sd_values)
```

```
## # A tibble: 4 x 4
##   Income_class Sd_Median_Age Sd_Youth_Unemployment_Rate Sd_Migration_Rate
##   <fct>          <dbl>          <dbl>          <dbl>
## 1 L              1.74             12.4             2.60
## 2 LM             5.08             11.3             2.21
## 3 UM             6.34             13.3            13.2
## 4 H              5.66             10.8             5.99
```

The standard deviation results do provide some additional insights. For the median age, we can see that the sd are relatively low across all income classes, meaning less variability, therefore similar age distributions. The sd of youth unemployment rate is higher compared to the median age. This indicates a bigger variability in the rates, suggesting that employment conditions vary more among countries. For the migration rate, the sd is lower for L, LM and H, and is the highest in UM class. This might show how the countries with UM income have more diverse migration patterns.

O

Repeat the analysis in m. for each income status and continent combination. Discuss the results.

```
table_of_average_values_continent <- df_vars %>%
  filter(!is.na(Income_class)) %>%
  mutate(Income_class = factor(Income_class, levels = c("L", "LM", "UM", "H"))) %>%
  group_by(Income_class, continent) %>%
  summarise(
    Avg_Median_Age = mean(Median_Age, na.rm = TRUE),
    Avg_Youth_Unemployment_Rate = mean(Youth_unemployment_rate, na.rm = TRUE),
    Avg_Migration_Rate = mean(Migration_rate, na.rm = TRUE),
    .groups = 'drop'
  )

print(table_of_average_values_continent)
```

```
## # A tibble: 20 x 5
##   Income_class continent Avg_Median_Age Avg_Youth_Unemployment_Rate
##   <fct>          <chr>          <dbl>          <dbl>
## 1 L            Africa              18.2             12.9
## 2 L            Asia               19.5             17.6
## 3 L            <NA>               19.9             13.1
## 4 LM          Africa              22.1             20.8
## 5 LM          Americas             25.4             10.2
## 6 LM          Asia               27.5             13.1
## 7 LM          Europe              41.2             17.9
## 8 LM          Oceania             24.4             13.5
## 9 UM          Africa              25.6             39.3
## 10 UM         Americas             30.9             15.3
## 11 UM         Asia               31.9             19.4
## 12 UM         Europe              40.3             19.7
## 13 UM         Oceania             26.3             15.7
```

```
## 14 UM          <NA>          36.9          44.6
## 15 H           Africa         36.8          11.6
## 16 H           Americas       38.3          19.1
## 17 H           Asia           34.4          11.1
## 18 H           Europe         43.2          15.7
## 19 H           Oceania        33           25.7
## 20 H           <NA>          35           25.8
## # i 1 more variable: Avg_Migration_Rate <dbl>
```

From the results we can observe the socioeconomic differences across the continents. We see that lower income countries have lower median ages, while higher income ones have older populations. Also, unemployment rates vary widely across regions, some having significantly higher rates than others. For L class, we have similar rates; for LM and UM, Africa has the highest rate; for H, Oceania has the highest rate. Furthermore, migration patterns also differ. We have regions experimenting negative migration rates, while others having positive ones. L class has only negative rates; for LM, Europe has positive rate; for UM Africa has the positive one, while Oceania followed by Asia have the lowest rates; for H, we see only Oceania with negative rate. Lastly, we have for each L, UM, H classes, an observation with NA values for the continent variable. This is important to note as it might affect the accuracy of the results.

p

Identify countries which are doing well in terms of both youth unemployment and net migration rate (in the top 25% of their respective continent in terms of net migration rate and in the bottom 25% of their respective continent in terms of youth unemployment).

```
top_countries <- df_vars %>%
  group_by(continent) %>%
  mutate(
    top_migration = quantile(Migration_rate, 0.75, na.rm = TRUE),
    bottom_unemployment = quantile(Youth_unemployment_rate, 0.25, na.rm = TRUE)
  ) %>%
  filter(Migration_rate >= top_migration & Youth_unemployment_rate <= bottom_unemployment) %>%
  select(Country, continent, Migration_rate, Youth_unemployment_rate)
```

```
top_countries
```

```
## # A tibble: 17 x 4
## # Groups:   continent [6]
##   Country          continent Migration_rate Youth_unemployment_rate
##   <chr>             <chr>          <dbl>          <dbl>
## 1 Czechia          <NA>            2.3            6.7
## 2 Switzerland      Europe           4.6            7.9
## 3 Malta            Europe           6.6            9.1
## 4 Macau            <NA>            3.3            5.3
## 5 Norway           Europe           4              9.7
## 6 United States     Americas         3              8.6
## 7 United Arab Emirates Asia             7.6            6.9
## 8 Palau            Oceania          0.9            5.6
## 9 Qatar            Asia             6.5            0.4
## 10 Bahrain         Asia            10.6            5.3
## 11 Kazakhstan      Asia             0.4            3.8
```

## 12 Israel	Asia	2.1	7.2
## 13 Papua New Guinea	Oceania	0	3.6
## 14 Madagascar	Africa	0	1
## 15 Togo	Africa	0	3.9
## 16 Guinea	Africa	0	1
## 17 Benin	Africa	0.3	5.6

The results show diverse regions. Countries like Qatar, Switzerland, United Arab Emirates etc. present favorable results for both rates. Meanwhile, Africa countries, which have lower unemployment rates, have also lower migration rates.

r

Export the final data set to a csv with “;” separator and “.” as a symbol for missing values; no rownames should be included in the csv. Upload the .csv to TUWEL together with your .Rmd and .html (or .pdf).

```
write.table(df_vars, file = "final_data.csv", sep = ";", na = ".", row.names = FALSE, col.names = TRUE)
```